

P8130_final_project

Leonor Rui

2024-12-03

Appendix

- Data Import

```
survival_df = read_csv("data/Project_2_data.csv") |>
  janitor::clean_names()
```

- Data Description

```
str(survival_df)
```

```
## spc_tbl_ [4,024 x 16] (S3: spec_tbl_df/tbl_df/tbl/data.frame)
## $ age          : num [1:4024] 68 50 58 58 47 51 51 40 40 69 ...
## $ race         : chr [1:4024] "White" "White" "White" "White" ...
## $ marital_status : chr [1:4024] "Married" "Married" "Divorced" "Married" ...
## $ t_stage      : chr [1:4024] "T1" "T2" "T3" "T1" ...
## $ n_stage      : chr [1:4024] "N1" "N2" "N3" "N1" ...
## $ x6th_stage   : chr [1:4024] "IIA" "IIIA" "IIIC" "IIA" ...
## $ differentiate : chr [1:4024] "Poorly differentiated" "Moderately differentiated" "Moderate" ...
## $ grade        : chr [1:4024] "3" "2" "2" "3" ...
## $ a_stage      : chr [1:4024] "Regional" "Regional" "Regional" "Regional" ...
## $ tumor_size   : num [1:4024] 4 35 63 18 41 20 8 30 103 32 ...
## $ estrogen_status : chr [1:4024] "Positive" "Positive" "Positive" "Positive" ...
## $ progesterone_status : chr [1:4024] "Positive" "Positive" "Positive" "Positive" ...
## $ regional_node_examined: num [1:4024] 24 14 14 2 3 18 11 9 20 21 ...
## $ reginol_node_positive : num [1:4024] 1 5 7 1 1 2 1 1 18 12 ...
## $ survival_months : num [1:4024] 60 62 75 84 50 89 54 14 70 92 ...
## $ status        : chr [1:4024] "Alive" "Alive" "Alive" "Alive" ...
## - attr(*, "spec")=
## .. cols(
## ..   Age = col_double(),
## ..   Race = col_character(),
## ..   'Marital Status' = col_character(),
## ..   'T Stage' = col_character(),
## ..   'N Stage' = col_character(),
## ..   '6th Stage' = col_character(),
## ..   differentiate = col_character(),
## ..   Grade = col_character(),
## ..   'A Stage' = col_character(),
```

```
## .. 'Tumor Size' = col_double(),
## .. 'Estrogen Status' = col_character(),
## .. 'Progesterone Status' = col_character(),
## .. 'Regional Node Examined' = col_double(),
## .. 'Reginol Node Positive' = col_double(),
## .. 'Survival Months' = col_double(),
## .. Status = col_character()
## .. )
## - attr(*, "problems")=<externalptr>
```

Numeric variables include age, tumor_size, regional_node_examined, reginol_node_positive, and survival_months.

These are continuous variables that can be used for our later regression analysis.

Categorical variables include race, marital_status, t_stage, n_stage, x6th_stage, differentiate, grade, a_stage, estrogen_status, progesterone_status, and status.

Then we will convert these variables into factors.

```
survival_df = survival_df |>
  mutate(
    race = factor(race),
    marital_status = factor(marital_status),
    t_stage = factor(t_stage),
    n_stage = factor(n_stage),
    x6th_stage = factor(x6th_stage),
    differentiate = factor(differentiate),
    grade = factor(grade),
    a_stage = factor(a_stage),
    estrogen_status = factor(estrogen_status),
    progesterone_status = factor(progesterone_status),
    status = factor(status)
  )
```

```
summary(survival_df)
```

```
##      age      race      marital_status t_stage  n_stage  x6th_stage
## Min.   :30.00 Black: 291 Divorced : 486 T1:1603 N1:2732 IIA :1305
## 1st Qu.:47.00 Other: 320 Married  :2643 T2:1786 N2: 820 IIB :1130
## Median :54.00 White:3413 Separated: 45 T3: 533 N3: 472 IIIA:1050
## Mean   :53.97          Single  : 615 T4: 102          IIIB: 67
## 3rd Qu.:61.00          Widowed  : 235          IIIC: 472
## Max.   :69.00
##
##      differentiate      grade      a_stage
## Moderately differentiated:2351 1 : 543 Distant : 92
## Poorly differentiated :1111 2 :2351 Regional:3932
## Undifferentiated : 19 3 :1111
## Well differentiated : 543 anaplastic; Grade IV: 19
##
##
##      tumor_size      estrogen_status progesterone_status regional_node_examined
## Min.   : 1.00 Negative: 269 Negative: 698 Min.   : 1.00
## 1st Qu.: 16.00 Positive:3755 Positive:3326 1st Qu.: 9.00
## Median : 25.00
```

```
## Mean      : 30.47                      Mean      :14.36
## 3rd Qu.: 38.00                      3rd Qu.:19.00
## Max.      :140.00                    Max.      :61.00
## reginol_node_positive survival_months status
## Min.      : 1.000      Min.      : 1.0    Alive:3408
## 1st Qu.: 1.000      1st Qu.: 56.0    Dead : 616
## Median : 2.000      Median : 73.0
## Mean      : 4.158      Mean      : 71.3
## 3rd Qu.: 5.000      3rd Qu.: 90.0
## Max.      :46.000      Max.      :107.0
```

The majority of patients in the dataset are White, accounting for approximately 84.82% of the total population. Black patients make up 7.23%, and patients classified as “Other” constitute 7.95%. This imbalance suggests that the dataset is heavily skewed towards White patients, which could influence the generalizability of the findings to other racial groups.

The wide range of values in variables such as `tumor_size`, `regional_node_examined`, and `survival_months` indicates the need to explore relationships and their potential nonlinearities with survival, giving us a possible analytical regression model.

```
colSums(is.na(survival_df))
```

```
##          age          race          marital_status
##          0          0          0
##          t_stage          n_stage          x6th_stage
##          0          0          0
##          differentiate          grade          a_stage
##          0          0          0
##          tumor_size          estrogen_status          progesterone_status
##          0          0          0
## regional_node_examined reginol_node_positive          survival_months
##          0          0          0
##          status
##          0
```

We can conclude that no missing values are present in this dataset across all variables.

```
survival_df |>
  group_by(differentiate, race) |>
  summarise(count = n(), .groups = "drop") |>
  pivot_wider(
    names_from = differentiate,
    values_from = count,
    values_fill = list(count = 0)
  )
```

```
## # A tibble: 3 x 5
##   race 'Moderately differentiated' 'Poorly differentiated' Undifferentiated
##   <fct>          <int>          <int>          <int>
## 1 Black          141          115           3
## 2 Other          180           94           0
## 3 White         2030          902          16
## # i 1 more variable: 'Well differentiated' <int>
```

This table shows the frequency of different levels of `differentiate` by races.

```
survival_df |>
  group_by(x6th_stage, status) |>
  summarise(count = n(), .groups = "drop") |>
  pivot_wider(
    names_from = status,
    values_from = count
  )
```

```
## # A tibble: 5 x 3
##   x6th_stage Alive  Dead
##   <fct>      <int> <int>
## 1 IIA         1209    96
## 2 IIB          995   135
## 3 IIIA         866   184
## 4 IIIB          47    20
## 5 IIIC         291   181
```


This table shows the frequency of different levels of `status` by 6th stage.

- Data Visualization

Distributions of the numeric variables

Distribution of age

```
survival_df |>
  ggplot(aes(age)) +
  geom_histogram(fill = "light blue", color = "black") +
  theme_minimal() +
  labs(
    title = "Distribution of age",
    x = "Age",
    y = "Frequency"
  )
```




The histogram shows the age distribution of patients. Most patients are aged between 40 and 70 years. The data is well spread across middle and older age groups, making it possible for age-related analysis. Therefore, age will likely be a significant predictor for later analysis.

Distribution of tumor size

```
ggplot(survival_df, aes(x = tumor_size)) +
  geom_histogram(fill = "light blue", color = "black") +
  scale_x_continuous(breaks = seq(0, max(survival_df$tumor_size, na.rm = TRUE), by = 5)) +
  labs(
    title = "Distribution of Tumor Size",
    x = "Tumor Size (mm)",
    y = "Frequency"
  )
```

Distribution of Tumor Size




This is the distribution of all tumor sizes, and most of the tumor sizes are smaller than 50 mm. We can find that the most frequent size is around 19 mm, followed by around 14 mm. This distribution is right-skewed, so we will use the log transformation for this variable.

Distribution of examined regional node

```
ggplot(survival_df, aes(x = regional_node_examined)) +
  geom_histogram(fill = "light blue", color = "black") +
  scale_x_continuous(breaks = seq(0, max(survival_df$regional_node_examined, na.rm = TRUE), by = 5)) +
  labs(
    title = "Distribution of Examined Regional Node",
    x = "Examined Regional Node",
    y = "Frequency"
  )
```

Distribution of Examined Regional Node




This plot maps the frequency of different number of examined regional nodes for each subject. The number of examined regional nodes for most subjects are smaller than 30, and the subjects with nearly 12 examined regional nodes are the most.

Distribution of positive regional node

```
ggplot(survival_df, aes(x = reginol_node_positive)) +
  geom_histogram(fill = "light blue", color = "black") +
  scale_x_continuous(breaks = seq(0, max(survival_df$reginol_node_positive, na.rm = TRUE), by = 5)) +
  labs(
    title = "Distribution of Positive Reginol Node",
    x = "Positive Reginol Node",
    y = "Frequency"
  )
```

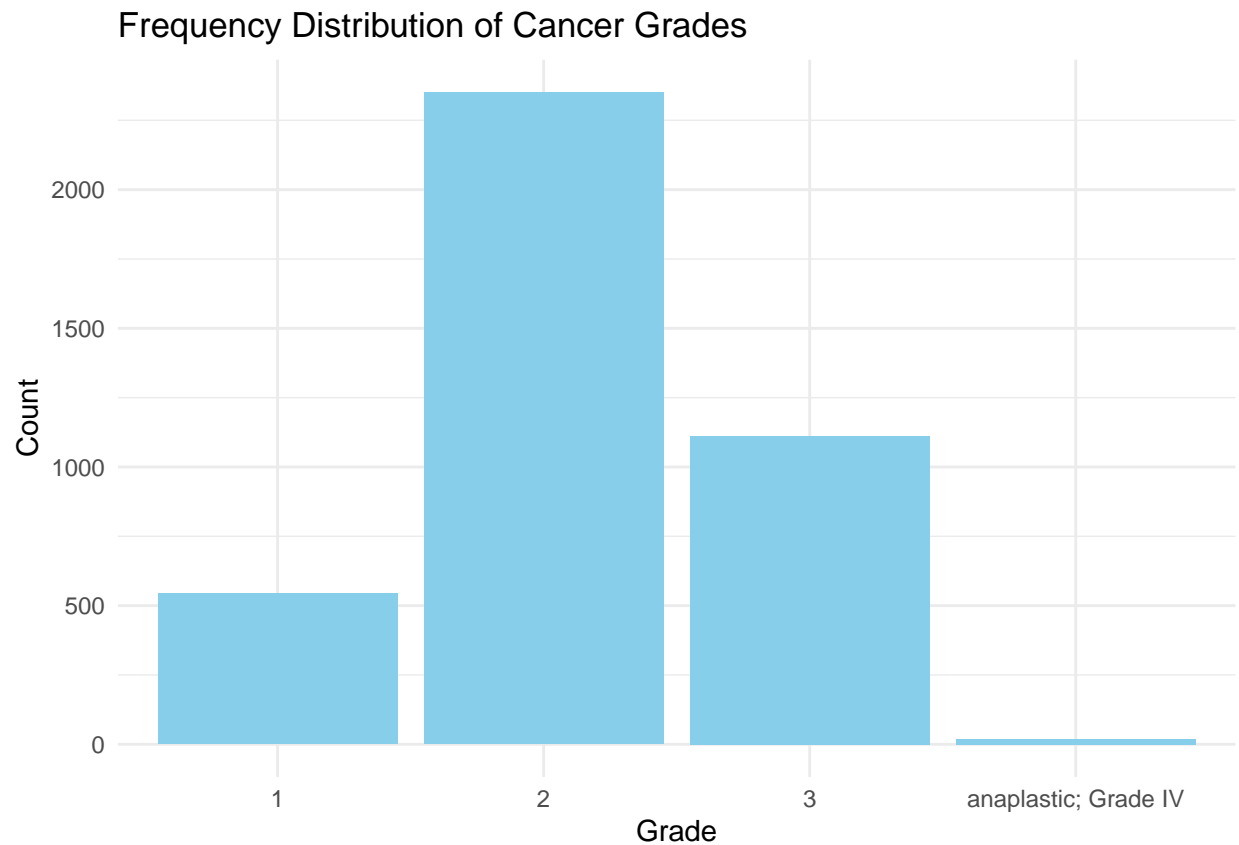
Distribution of Positive Reginol Node



Then is the distribution of different number of positive reginol node for each subject. Over 2500 subjects only have 1 or 2 positive reginol nodes, which is the most frequent number of positive reginol nodes. It is strongly right-skewed, so we will use the log transformation for this variable.

Distribution of Cancer Grades

```
ggplot(survival_df, aes(x = grade)) +  
  geom_bar(fill = "skyblue") +  
  labs(title = "Frequency Distribution of Cancer Grades",  
        x = "Grade",  
        y = "Count") +  
  theme_minimal()
```





This bar chart provides an overview of how cancer cases are distributed across grades. Grade 2 represents the majority of cases, suggesting it is the most frequently observed grade, while Grade IV is exceedingly rare.

Between Variables

Distribution of survival months by status


```
ggplot(survival_df, aes(x = survival_months, fill = status)) +
  geom_histogram(binwidth = 5, position = "dodge") +
  labs(title = "Distribution of Survival Months", x = "Survival Months", y = "Frequency") +
  theme_minimal()
```



The Dead group is concentrated in the shorter survival months, while the Alive group is predominant in longer survival months, particularly beyond 60 months.

The distribution of the tumor sizes by t_stage

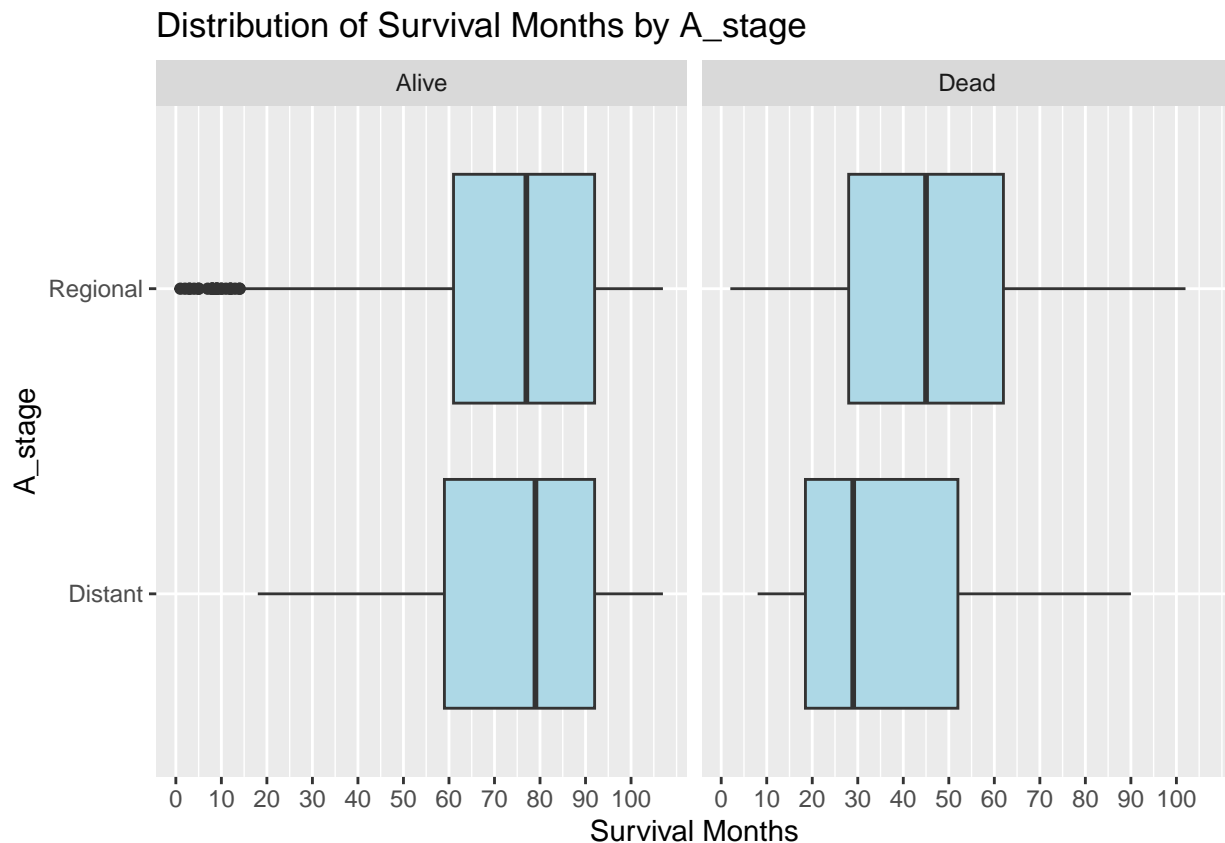
```
ggplot(survival_df, aes(x = tumor_size, y = t_stage)) +
  geom_boxplot(fill = "light blue") +
  scale_x_continuous(breaks = seq(0, max(survival_df$tumor_size, na.rm = TRUE), by = 10)) +
  labs(
    title = "Distribution of Tumor Sizes by T_stage",
    x = "Tumor Sizes",
    y = "T_stage"
  )
```



In this plot, we explore the tumor size distribution at different T stages. From T1 to T3, as the stage changes, both the mean tumor sizes and IQR become larger. At T4 stage, the IQR of tumor sizes is much larger than others, and the mean size is smaller than the mean size at T3 stage. There are some outliers both at T1 stage and T3 stage.

The distribution of survival months by a_stage based on status(alive/dead)

```
ggplot(survival_df, aes(x = survival_months, y = a_stage)) +
  geom_boxplot(fill = "light blue") +
  scale_x_continuous(breaks = seq(0, max(survival_df$survival_months, na.rm = TRUE), by = 10)) +
  labs(
    title = "Distribution of Survival Months by A_stage",
    x = "Survival Months",
    y = "A_stage"
  ) +
  facet_grid(~ status)
```

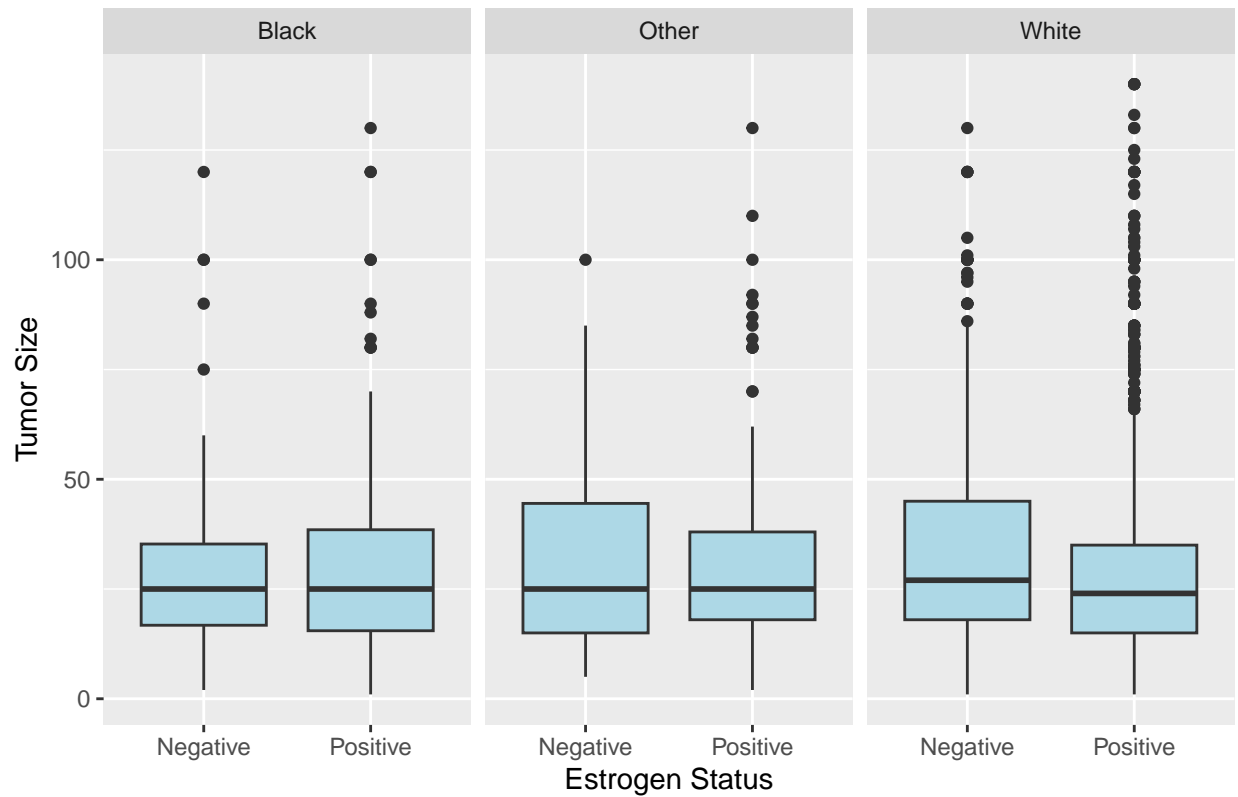


Through this plot, we can find that subjects with Distant stage have fewer survival months than subjects with Regional stage. However, the IQR of the survival months of subjects with Distant stage is much larger than subjects with Regional stage.

Distribution of Estrogen Status by Tumor Size Based on race

```
ggplot(survival_df, aes(x = progesterone_status, y = tumor_size)) +
  geom_boxplot(fill = "light blue") +
  labs(
    title = "Distribution of Estrogen Status by Tumor Size",
    x = "Estrogen Status",
    y = "Tumor Size"
  ) +
  facet_grid(~ race)
```


Distribution of Estrogen Status by Tumor Size



While the overall patterns are consistent, with Negative estrogen status generally associated with slightly larger tumor sizes, the variability and prevalence of outliers differ between groups. The White group shows the greatest spread in tumor size, while the Other group displays the least variability.

Tumor Size Distribution by Grade

```
ggplot(survival_df, aes(x = grade, y = tumor_size)) +  
  geom_boxplot(fill = "lightblue") +  
  labs(title = "Tumor Size Distribution by Grade",  
        x = "Grade",  
        y = "Tumor Size (mm)") +  
  theme_minimal()
```




Lower grades (1–3) exhibit comparable tumor size distributions, with slight increases in variability as the grade increases.

Grade IV stands out due to its higher median and broader range, suggesting that more aggressive tumor grades are associated with larger tumor sizes.

Proportion of Grade by Estrogen Status


```
ggplot(survival_df, aes(x = grade, fill = estrogen_status)) +
  geom_bar(position = "fill") +
  labs(title = "Proportion of Grade by Estrogen Status",
       x = "Grade",
       y = "Proportion",
       fill = "Race") +
  theme_minimal()
```



As tumor grade increases, the proportion of Negative estrogen status gradually increases, becoming more prominent in the anaplastic Grade IV category. Conversely, the dominance of the Positive estrogen status decreases with higher tumor grades.

Proportion of Grade by Marital Status

```
ggplot(survival_df, aes(x = grade, fill = marital_status)) +
  geom_bar(position = "fill") +
  labs(title = "Proportion of Grade by Marital Status",
       x = "Grade",
       y = "Proportion",
       fill = "Race") +
  theme_minimal()
```



Across all grades, the “Married” group consistently constitutes the largest proportion of individuals, dominating every tumor grade category.


The “Single” group is the second-largest proportion in most grades, particularly Grades 2 and 3.

The “Widowed” group and “Divorced” group make up smaller proportions across all tumor grades.

Relationship Between Age and Tumor Size across status

```
ggplot(survival_df, aes(x = age, y = tumor_size, color = status)) +
  geom_point(alpha = 0.5, size = 0.6) +
  geom_smooth(method = "lm", color = "green") +
  facet_wrap(~ status) +
  labs(title = "Relationship Between Age and Tumor Size",
       x = "Age (years)",
       y = "Tumor Size (mm)",
       color = "Status") +
  theme_minimal()
```


Relationship Between Age and Tumor Size




This figure highlights the differences in tumor size distribution and trends with age between individuals who are alive and those who are deceased. While the “Alive” group shows no significant relationship between age and tumor size, the “Dead” group exhibits a pattern where larger tumors are associated with younger ages.

Age vs. Tumor Size Across Grades

```
ggplot(survival_df, aes(x = age, y = tumor_size)) +  
  geom_point(color = "lightblue", size = 0.8, alpha = 0.5) +  
  geom_smooth(method = "lm") +  
  facet_wrap(~ marital_status) +  
  labs(title = "Age vs. Tumor Size Across Grades",  
       x = "Age (years)",  
       y = "Tumor Size (mm)") +  
  theme_minimal()
```

Age vs. Tumor Size Across Grades



Divorced: Tumor size seems to remain fairly constant with age, as the trend line is relatively flat.

Married: A slight negative trend is observable, suggesting that tumor size may decrease marginally with age.

Separated: The data is sparse, but the trend shows a slightly negative relationship, with wide confidence intervals due to fewer observations.

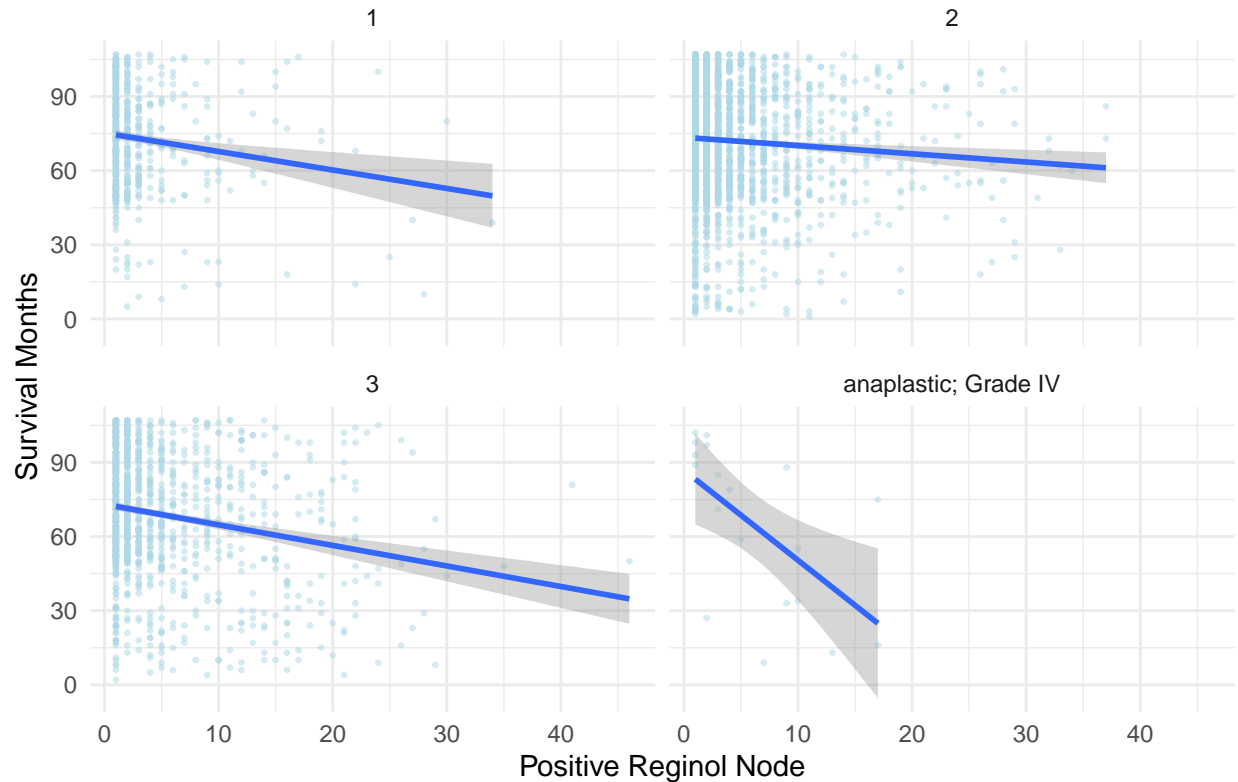
Single: A modest negative trend is observed, indicating a potential decline in tumor size with increasing age.

Widowed: A more apparent negative trend is evident compared to other groups, suggesting a stronger decrease in tumor size with age.

Positive Regiol Node vs Survival Months Across Cancer Grade

```
ggplot(survival_df, aes(x = regiol_node_positive, y = survival_months)) +
  geom_point(color = "light blue", size = 0.5, alpha = 0.5) +
  facet_wrap(~grade) +
  geom_smooth(method = "lm") +
  labs(
    title = "Distribution of Positive Regiol Node and Survival Months by Cancer Grade",
    x = "Positive Regiol Node",
    y = "Survival Months"
  ) +
  theme_minimal()
```

Distribution of Positive Reginol Node and Survival Months by Cancer Grade



According to the trend lines, as the cancer grade increases, the negative correlation between the number of positive reginol nodes and the survival months becomes stronger. At the Grade IV, the correlation is strong. As the number of positive reginol nodes increases, the survival months will decrease.


Transformations

```
survival_df = survival_df |>
  mutate(
    log_tumor_size = log(tumor_size),
    log_reginol_node_positive = log(reginol_node_positive)
  )
```

Since variables `tumor_size` and `reginol_node_positive` are skewed to the right, we need to use the log transformation and add new variables `log_tumor_size` and `log_reginol_node_positive` for further analysis.

```
survival_df |>
  pivot_longer(
    cols = c(age, tumor_size, regional_node_examined, reginol_node_positive, survival_months),
    names_to = "variable",
    values_to = "value"
  ) |>
  ggplot(aes(x = value)) +
```

```
geom_histogram() +  
facet_wrap(variable ~ ., scales = "free")
```



Model Building

Logistic Regression Model


Adjusted R^2 increases as predictor number increases, but need to simplify the model.

```
subsets = regsubsets(status ~ age + race + marital_status + t_stage + n_stage + x6th_stage + differentia
```

Reordering variables and trying again:

```
subset_summary = summary(subsets)  
  
plot(subset_summary$adjr2, type = "b",  
     xlab = "Number of Predictors", ylab = "Adjusted R^2",  
     main = "Distribution of Adjusted R^2 of Different Models")
```

Distribution of Adjusted R² of Different Models



Automated procedure

```
full_glm = glm(status ~ age + race + marital_status + t_stage + n_stage + x6th_stage + differentiate + grade + a_stage + tumor_size + estrogen_status + progesterone_status + regional_node_examined + reginol_node_positive)
step(full_glm)
```

```
## Start: AIC=3002
## status ~ age + race + marital_status + t_stage + n_stage + x6th_stage +
##   differentiate + grade + a_stage + tumor_size + estrogen_status +
##   progesterone_status + regional_node_examined + reginol_node_positive
##
## Step: AIC=3002
## status ~ age + race + marital_status + t_stage + n_stage + x6th_stage +
##   differentiate + a_stage + tumor_size + estrogen_status +
##   progesterone_status + regional_node_examined + reginol_node_positive
##
##           Df Deviance    AIC
## - x6th_stage      3   2953.7 2997.7
## - tumor_size      1   2952.0 3000.0
## - a_stage         1   2952.0 3000.0
## - t_stage         3   2957.2 3001.2
## <none>            2952.0 3002.0
## - marital_status  4   2960.1 3002.1
## - n_stage         1   2958.9 3006.9
## - race            2   2967.1 3013.1
```

```

## - estrogen_status      1  2969.2 3017.2
## - age                  1  2970.8 3018.8
## - progesterone_status  1  2972.0 3020.0
## - regional_node_examined 1  2978.7 3026.7
## - reginol_node_positive 1  2978.8 3026.8
## - differentiate        3  2986.0 3030.0
##
## Step: AIC=2997.73
## status ~ age + race + marital_status + t_stage + n_stage + differentiate +
##      a_stage + tumor_size + estrogen_status + progesterone_status +
##      regional_node_examined + reginol_node_positive
##
##              Df Deviance    AIC
## - tumor_size      1  2953.8 2995.8
## - a_stage          1  2953.8 2995.8
## <none>              2953.7 2997.7
## - marital_status   4  2961.7 2997.7
## - n_stage          2  2965.5 3005.5
## - t_stage          3  2968.8 3006.8
## - race             2  2969.2 3009.2
## - estrogen_status  1  2971.2 3013.2
## - age              1  2972.1 3014.1
## - progesterone_status 1  2973.6 3015.6
## - regional_node_examined 1  2980.4 3022.4
## - reginol_node_positive 1  2981.1 3023.1
## - differentiate    3  2987.4 3025.4
##
## Step: AIC=2995.75
## status ~ age + race + marital_status + t_stage + n_stage + differentiate +
##      a_stage + estrogen_status + progesterone_status + regional_node_examined +
##      reginol_node_positive
##
##              Df Deviance    AIC
## - a_stage          1  2953.8 2993.8
## <none>              2953.8 2995.8
## - marital_status   4  2961.8 2995.8
## - n_stage          2  2965.7 3003.7
## - race             2  2969.2 3007.2
## - estrogen_status  1  2971.2 3011.2
## - age              1  2972.1 3012.1
## - progesterone_status 1  2973.7 3013.7
## - t_stage          3  2980.8 3016.8
## - regional_node_examined 1  2980.6 3020.6
## - reginol_node_positive 1  2981.1 3021.1
## - differentiate    3  2987.4 3023.4
##
## Step: AIC=2993.77
## status ~ age + race + marital_status + t_stage + n_stage + differentiate +
##      estrogen_status + progesterone_status + regional_node_examined +
##      reginol_node_positive
##
##              Df Deviance    AIC
## <none>              2953.8 2993.8
## - marital_status   4  2961.8 2993.8

```

```
## - n_stage          2    2965.8 3001.8
## - race             2    2969.3 3005.3
## - estrogen_status  1    2971.3 3009.3
## - age              1    2972.1 3010.1
## - progesterone_status 1    2973.7 3011.7
## - t_stage          3    2982.7 3016.7
## - regional_node_examined 1    2980.6 3018.6
## - reginol_node_positive 1    2981.1 3019.1
## - differentiate     3    2987.4 3021.4

##
## Call:  glm(formula = status ~ age + race + marital_status + t_stage +
##          n_stage + differentiate + estrogen_status + progesterone_status +
##          regional_node_examined + reginol_node_positive, family = binomial,
##          data = survival_df)
##
## Coefficients:
##              (Intercept)                  age
##                -1.75107                0.02380
##                raceOther                raceWhite
##                -0.93458                -0.51482
##      marital_statusMarried      marital_statusSeparated
##                -0.21103                0.66914
##      marital_statusSingle      marital_statusWidowed
##                -0.06461                0.01749
##                t_stageT2                t_stageT3
##                0.41111                0.55159
##                t_stageT4                n_stageN2
##                1.09878                0.43629
##                n_stageN3 differentiatePoorly differentiated
##                0.58717                0.38627
##      differentiateUndifferentiated      differentiateWell differentiated
##                1.33210                -0.53276
##      estrogen_statusPositive      progesterone_statusPositive
##                -0.74802                -0.58416
##      regional_node_examined      reginol_node_positive
##                -0.03594                0.07968
##
## Degrees of Freedom: 4023 Total (i.e. Null);  4004 Residual
## Null Deviance:      3445
## Residual Deviance: 2954  AIC: 2994
```

```
automated_glm = glm(formula = status ~ age + race + marital_status + t_stage + n_stage + differentiate +
summary(automated_glm)
```

```
##
## Call:
## glm(formula = status ~ age + race + marital_status + t_stage +
##      n_stage + differentiate + estrogen_status + progesterone_status +
##      regional_node_examined + reginol_node_positive, family = binomial,
##      data = survival_df)
##
```

```
## Coefficients:
##               Estimate Std. Error z value Pr(>|z|)
## (Intercept)    -1.751071   0.405332  -4.320 1.56e-05 ***
## age              0.023801   0.005610   4.243 2.21e-05 ***
## raceOther       -0.934582   0.248456  -3.762 0.000169 ***
## raceWhite       -0.514824   0.161664  -3.185 0.001450 **
## marital_statusMarried -0.211035   0.141636  -1.490 0.136229
## marital_statusSeparated  0.669140   0.388139   1.724 0.084713 .
## marital_statusSingle -0.064613   0.174816  -0.370 0.711675
## marital_statusWidowed  0.017493   0.221070   0.079 0.936929
## t_stageT2        0.411113   0.113030   3.637 0.000276 ***
## t_stageT3        0.551595   0.148772   3.708 0.000209 ***
## t_stageT4        1.098780   0.244529   4.493 7.01e-06 ***
## n_stageN2        0.436292   0.128371   3.399 0.000677 ***
## n_stageN3        0.587168   0.234546   2.503 0.012300 *
## differentiatePoorly differentiated  0.386267   0.104812   3.685 0.000228 ***
## differentiateUndifferentiated  1.332103   0.530034   2.513 0.011963 *
## differentiateWell differentiated -0.532760   0.183777  -2.899 0.003744 **
## estrogen_statusPositive -0.748017   0.177506  -4.214 2.51e-05 ***
## progesterone_statusPositive -0.584158   0.127516  -4.581 4.63e-06 ***
## regional_node_examined -0.035938   0.007172  -5.011 5.41e-07 ***
## reginol_node_positive  0.079684   0.015301   5.208 1.91e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##    Null deviance: 3444.7  on 4023  degrees of freedom
## Residual deviance: 2953.8  on 4004  degrees of freedom
## AIC: 2993.8
##
## Number of Fisher Scoring iterations: 5
```

```
vif(automated_glm)
```


```
##               GVIF Df GVIF^(1/(2*Df))
## age              1.107154  1      1.052214
## race             1.062865  2      1.015359
## marital_status    1.129117  4      1.015295
## t_stage           1.101923  3      1.016308
## n_stage           3.806780  2      1.396817
## differentiate     1.117097  3      1.018627
## estrogen_status   1.475424  1      1.214670
## progesterone_status 1.427536  1      1.194795
## regional_node_examined 1.477809  1      1.215651
## reginol_node_positive 4.248357  1      2.061154
```

Model Diagnostics


```
augment(automated_glm) |>
  ggplot(aes(x = .fitted, y = .std.resid)) +
```



```
geom_point() +  
geom_smooth(se = FALSE) +  
labs(x = "Fitted value", y = "Residual")
```




```
augment_quantile(automated_glm) |>  
ggplot(aes(x = .fitted, y = .quantile.resid)) +  
geom_point() +  
geom_smooth(se = FALSE) +  
labs(x = "Fitted value", y = "Randomized quantile residual")
```



By randomizing the quantile residuals, we resolve the problem that the RVF plot always shows a pattern in logistic regression because of the binary response variable. Since in the randomized quantile residual vs. fitted value plot, the residuals distribute randomly around the 0.5 horizontal line, the residual assumption is met and the model is a good fit.

```
plot(automated_glm, which = 5)
```



`glm(status ~ age + race + marital_status + t_stage + n_stage + differentiat ...`

The residual vs. leverage plot indicates that observations 3527, 1561, and 3074 may be potential outliers, but they are not necessarily influential.

Likelihood ratio test

Test the validity of the automated logistic regression model

```
lrtest(automated_glm, full_glm)
```

```
## Likelihood ratio test
##
## Model 1: status ~ age + race + marital_status + t_stage + n_stage + differentiate +
##   estrogen_status + progesterone_status + regional_node_examined +
##   reginol_node_positive
## Model 2: status ~ age + race + marital_status + t_stage + n_stage + x6th_stage +
##   differentiate + grade + a_stage + tumor_size + estrogen_status +
##   progesterone_status + regional_node_examined + reginol_node_positive
##   #Df  LogLik Df  Chisq Pr(>Chisq)
## 1   20 -1476.9
## 2   25 -1476.0  5  1.7712    0.8798
```

```
final_glm = automated_glm
```

According to the likelihood ratio test, the full model is not significantly more effective in predicting survival status than the automated model, so we will use the automated model as the final logistic model.

Odds Ratios

```
final_glm_summary = summary(final_glm)

final_glm_df = as.data.frame(final_glm_summary$coefficients) |>
  janitor::clean_names() |>
  mutate(
    odds_ratio = exp(estimate)
  ) |>
  rename(p_value = pr_z)

final_glm_df
```

| | estimate | std_error | z_value |
|---------------------------------------|--------------|-------------|-------------|
| ## (Intercept) | -1.75107145 | 0.405332256 | -4.32008906 |
| ## age | 0.02380073 | 0.005609963 | 4.24258193 |
| ## raceOther | -0.93458158 | 0.248455746 | -3.76156155 |
| ## raceWhite | -0.51482432 | 0.161663812 | -3.18453657 |
| ## marital_statusMarried | -0.21103457 | 0.141635616 | -1.48998234 |
| ## marital_statusSeparated | 0.66913999 | 0.388138549 | 1.72397201 |
| ## marital_statusSingle | -0.06461326 | 0.174815803 | -0.36960768 |
| ## marital_statusWidowed | 0.01749330 | 0.221070476 | 0.07912998 |
| ## t_stageT2 | 0.41111286 | 0.113029538 | 3.63721613 |
| ## t_stageT3 | 0.55159473 | 0.148771968 | 3.70765235 |
| ## t_stageT4 | 1.09878030 | 0.244529438 | 4.49344792 |
| ## n_stageN2 | 0.43629245 | 0.128370873 | 3.39868728 |
| ## n_stageN3 | 0.58716755 | 0.234545523 | 2.50342683 |
| ## differentiatePoorly differentiated | 0.38626692 | 0.104812369 | 3.68531809 |
| ## differentiateUndifferentiated | 1.33210322 | 0.530033961 | 2.51324127 |
| ## differentiateWell differentiated | -0.53276008 | 0.183776817 | -2.89895151 |
| ## estrogen_statusPositive | -0.74801661 | 0.177506343 | -4.21402748 |
| ## progesterone_statusPositive | -0.58415843 | 0.127515990 | -4.58106024 |
| ## regional_node_examined | -0.03593764 | 0.007171744 | -5.01100482 |
| ## reginol_node_positive | 0.07968369 | 0.015301426 | 5.20759867 |
| | p_value | odds_ratio | |
| ## (Intercept) | 1.559662e-05 | 0.1735879 | |
| ## age | 2.209628e-05 | 1.0240862 | |
| ## raceOther | 1.688559e-04 | 0.3927502 | |
| ## raceWhite | 1.449860e-03 | 0.5976056 | |
| ## marital_statusMarried | 1.362289e-01 | 0.8097461 | |
| ## marital_statusSeparated | 8.471290e-02 | 1.9525574 | |
| ## marital_statusSingle | 7.116748e-01 | 0.9374299 | |
| ## marital_statusWidowed | 9.369292e-01 | 1.0176472 | |
| ## t_stageT2 | 2.756007e-04 | 1.5084956 | |
| ## t_stageT3 | 2.091896e-04 | 1.7360193 | |
| ## t_stageT4 | 7.007917e-06 | 3.0005041 | |
| ## n_stageN2 | 6.771009e-04 | 1.5469611 | |
| ## n_stageN3 | 1.229971e-02 | 1.7988859 | |
| ## differentiatePoorly differentiated | 2.284172e-04 | 1.4714774 | |
| ## differentiateUndifferentiated | 1.196275e-02 | 3.7890041 | |
| ## differentiateWell differentiated | 3.744128e-03 | 0.5869826 | |
| ## estrogen_statusPositive | 2.508567e-05 | 0.4733044 | |
| ## progesterone_statusPositive | 4.626245e-06 | 0.5575749 | |

```
## regional_node_examined          5.414656e-07  0.9647004
## reginol_node_positive           1.913001e-07  1.0829445
```

Cross Validation

```
log_loss = function(actual, predicted) {
  -mean(actual * log(predicted) + (1 - actual) * log(1 - predicted))
}

survival_df = survival_df |>
  mutate(status = if_else(status == "dead", 0, 1))

cv_df =
  crossv_kfold(survival_df, k = 5) |>
  mutate(
    train = map(train, as_tibble),
    test = map(test, as_tibble)
  )

cv_res_df =
  cv_df |>
  mutate(
    final_model = map(train, \(x) glm(formula = status ~ age + race + marital_status + t_stage + n_stage,
                                     data = x, family = "binomial"))
    log_loss = map2_dbl(final_model, test, \(model, test_data) {
      predicted_probs = predict(model, newdata = test_data, type = "response")
      actual_outcomes = test_data$status
      log_loss(actual_outcomes, predicted_probs)
    })
  )

CV_log_loss = mean(cv_res_df$log_loss)
CV_log_loss
```

```
## [1] 2.90068e-12
```


Since the CV_log_loss is very small, the model is a good fit by cross validation.

- Evaluate the performance of your model(s). Is your model achieving similar performance between the majority race group “White” and the minority “Black” (or “Black” + “Other”)? If not, could you try to improve the fairness (i.e., reducing the gap of prediction performance between the majority and minority) of your model(s)?

Survival Analysis

```
km_fit = survfit(Surv(survival_months, status) ~ 1, data = survival_df)
ggsurvplot(km_fit, conf.int = TRUE, title = "Kaplan-Meier Survival Curve")
```

Kaplan–Meier Survival Curve



Fit a Cox Proportional Hazards Model

```
cox_model <- coxph(Surv(survival_months, status) ~ age + race + marital_status + t_stage + n_stage + x6th_stage + differentiate + grade + a_stage + tumor_size + estrogen_status + progesterone_status + regional_node_examined + reginol_node_positive, data = survival_df)
summary(cox_model)
```

Call:

```
## coxph(formula = Surv(survival_months, status) ~ age + race +
## marital_status + t_stage + n_stage + x6th_stage + differentiate +
## grade + a_stage + tumor_size + estrogen_status + progesterone_status +
## regional_node_examined + reginol_node_positive, data = survival_df)
##
```

```
## n= 4024, number of events= 4024
##
```

| | coef | exp(coef) | se(coef) | z |
|----------------------------|-----------|-----------|----------|--------|
| ## age | 0.003007 | 1.003011 | 0.001892 | 1.589 |
| ## raceOther | -0.230441 | 0.794183 | 0.082230 | -2.802 |
| ## raceWhite | -0.153526 | 0.857679 | 0.062208 | -2.468 |
| ## marital_statusMarried | 0.010499 | 1.010555 | 0.049801 | 0.211 |
| ## marital_statusSeparated | 0.420017 | 1.521987 | 0.157033 | 2.675 |
| ## marital_statusSingle | 0.006696 | 1.006719 | 0.061210 | 0.109 |
| ## marital_statusWidowed | 0.026505 | 1.026859 | 0.080713 | 0.328 |
| ## t_stageT2 | 0.056357 | 1.057975 | 0.075587 | 0.746 |
| ## t_stageT3 | 0.019494 | 1.019686 | 0.124698 | 0.156 |
| ## t_stageT4 | 0.130027 | 1.138859 | 0.204518 | 0.636 |
| ## n_stageN2 | -0.029644 | 0.970791 | 0.088322 | -0.336 |
| ## n_stageN3 | 0.049587 | 1.050837 | 0.118312 | 0.419 |
| ## x6th_stageIIB | -0.019870 | 0.980326 | 0.081350 | -0.244 |

```

## x6th_stageIIIA      0.042864  1.043796  0.104813  0.409
## x6th_stageIIIB     -0.159157  0.852862  0.229391 -0.694
## x6th_stageIIIC      NA          NA  0.000000    NA
## differentiatePoorly differentiated  0.020753  1.020969  0.037630  0.551
## differentiateUndifferentiated      0.072524  1.075218  0.232225  0.312
## differentiateWell differentiated   0.018197  1.018363  0.048089  0.378
## grade2              NA          NA  0.000000    NA
## grade3              NA          NA  0.000000    NA
## gradeanaplastic; Grade IV         NA          NA  0.000000    NA
## a_stageRegional     -0.004486  0.995524  0.119056 -0.038
## tumor_size          0.001543  1.001544  0.001597  0.966
## estrogen_statusPositive -0.165805  0.847211  0.073421 -2.258
## progesterone_statusPositive  0.046092  1.047171  0.048074  0.959
## regional_node_examined -0.002080  0.997922  0.002190 -0.950
## reginol_node_positive   0.013142  1.013228  0.006448  2.038
## Pr(>|z|)
## age                 0.11211
## raceOther           0.00507 **
## raceWhite           0.01359 *
## marital_statusMarried 0.83302
## marital_statusSeparated 0.00748 **
## marital_statusSingle  0.91289
## marital_statusWidowed 0.74262
## t_stageT2            0.45592
## t_stageT3            0.87577
## t_stageT4            0.52492
## n_stageN2            0.73714
## n_stageN3            0.67513
## x6th_stageIIB        0.80704
## x6th_stageIIIA       0.68257
## x6th_stageIIIB       0.48779
## x6th_stageIIIC       NA
## differentiatePoorly differentiated 0.58130
## differentiateUndifferentiated      0.75481
## differentiateWell differentiated   0.70513
## grade2              NA
## grade3              NA
## gradeanaplastic; Grade IV         NA
## a_stageRegional     0.96994
## tumor_size          0.33388
## estrogen_statusPositive 0.02393 *
## progesterone_statusPositive 0.33768
## regional_node_examined 0.34224
## reginol_node_positive  0.04155 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## exp(coef) exp(-coef) lower .95 upper .95
## age      1.0030      0.9970      0.9993      1.0067
## raceOther 0.7942      1.2592      0.6760      0.9331
## raceWhite 0.8577      1.1659      0.7592      0.9689
## marital_statusMarried 1.0106      0.9896      0.9166      1.1142
## marital_statusSeparated 1.5220      0.6570      1.1188      2.0705
## marital_statusSingle  1.0067      0.9933      0.8929      1.1350

```

```
## marital_statusWidowed          1.0269      0.9738      0.8766      1.2029
## t_stageT2                      1.0580      0.9452      0.9123      1.2269
## t_stageT3                      1.0197      0.9807      0.7986      1.3020
## t_stageT4                      1.1389      0.8781      0.7628      1.7004
## n_stageN2                      0.9708      1.0301      0.8165      1.1543
## n_stageN3                      1.0508      0.9516      0.8334      1.3251
## x6th_stageIIB                 0.9803      1.0201      0.8358      1.1498
## x6th_stageIIIA                1.0438      0.9580      0.8500      1.2818
## x6th_stageIIIB                0.8529      1.1725      0.5440      1.3370
## x6th_stageIIIC                NA          NA          NA          NA
## differentiatePoorly differentiated 1.0210      0.9795      0.9484      1.0991
## differentiateUndifferentiated    1.0752      0.9300      0.6821      1.6950
## differentiateWell differentiated  1.0184      0.9820      0.9268      1.1190
## grade2                        NA          NA          NA          NA
## grade3                        NA          NA          NA          NA
## gradeanaplastic; Grade IV      NA          NA          NA          NA
## a_stageRegional               0.9955      1.0045      0.7883      1.2572
## tumor_size                    1.0015      0.9985      0.9984      1.0047
## estrogen_statusPositive        0.8472      1.1803      0.7337      0.9783
## progesterone_statusPositive    1.0472      0.9550      0.9530      1.1506
## regional_node_examined         0.9979      1.0021      0.9936      1.0022
## reginol_node_positive          1.0132      0.9869      1.0005      1.0261
##
## Concordance= 0.545 (se = 0.005 )
## Likelihood ratio test= 62.84 on 24 df, p=3e-05
## Wald test = 66.31 on 24 df, p=8e-06
## Score (logrank) test = 66.66 on 24 df, p=7e-06
```

```
# Check proportional hazards assumption
cox.zph(cox_model)
```

```
##               chisq df      p
## age           0.435  1 0.50931
## race          4.621  2 0.09921
## marital_status 5.246  4 0.26293
## t_stage       12.768  3 0.00517
## n_stage       35.271  2 2.2e-08
## x6th_stage     3.005  3 0.39087
## differentiate  17.718  3 0.00050
## a_stage       18.098  1 2.1e-05
## tumor_size    12.283  1 0.00046
## estrogen_status 50.707  1 1.1e-12
## progesterone_status 43.882  1 3.5e-11
## regional_node_examined 0.147  1 0.70145
## reginol_node_positive 27.633  1 1.5e-07
## GLOBAL       124.982 24 1.2e-15
```

```
# Visualize adjusted survival curves
ggsurvplot(survfit(cox_model), data = survival_df, title = "Cox Model Survival Curve")
```


Cox Model Survival Curve

